

Lab 3

1. Use `fdisk -l` to locate information about the partition sizes.

```
[root@server ~]# fdisk -l
Disk /dev/nvme0n1: 50 GiB, 53687091200 bytes, 104857600 sectors
Disk model: VMware Virtual NVMe Disk
Units: sectors of 1 * 512 = 512 bytes
Sector size (logical/physical): 512 bytes / 512 bytes
I/O size (minimum/optimal): 512 bytes / 512 bytes
Disklabel type: dos
Disk identifier: 0x3395c45a

Device            Boot    Start        End    Sectors  Size Id Type
/dev/nvme0n1p1    *                2048    2099199    2097152    1G 83 Linux
/dev/nvme0n1p2                2099200  104857599  102758400    49G 8e Linux LVM

Disk /dev/nvme0n2: 50 GiB, 53687091200 bytes, 104857600 sectors
Disk model: VMware Virtual NVMe Disk
Units: sectors of 1 * 512 = 512 bytes
Sector size (logical/physical): 512 bytes / 512 bytes
I/O size (minimum/optimal): 512 bytes / 512 bytes

Disk /dev/mapper/rhel_server-root: 46.98 GiB, 50444894208 bytes, 98525184 sectors
Units: sectors of 1 * 512 = 512 bytes
Sector size (logical/physical): 512 bytes / 512 bytes
I/O size (minimum/optimal): 512 bytes / 512 bytes

Disk /dev/mapper/rhel_server-swap: 2.02 GiB, 2164260864 bytes, 4227072 sectors
Units: sectors of 1 * 512 = 512 bytes
Sector size (logical/physical): 512 bytes / 512 bytes
I/O size (minimum/optimal): 512 bytes / 512 bytes
[root@server ~]#
```

2. Use `fdisk` to add a new logical partition
that is 2GB in size. `Fdisk /dev/sdb`

```
[root@server ~]# fdisk /dev/nvme0n2
```

```
Welcome to fdisk (util-linux 2.37.4).
```

```
Changes will remain in memory only, until you decide  
to write them.
```

```
Be careful before using the write command.
```

```
Device does not contain a recognized partition table.  
Created a new DOS disklabel with disk identifier 0xc9  
bfbc6f.
```

```
Command (m for help): n
```

```
Partition type
```

```
  p   primary (0 primary, 0 extended, 4 free)
```

```
  e   extended (container for logical partitions)
```

```
Select (default p): p
```

```
Partition number (1-4, default 1): 1
```

```
First sector (2048-104857599, default 2048): +2G
```

```
Value out of range.
```

```
First sector (2048-104857599, default 2048):
```

```
Last sector, +/-sectors or +/-size{K,M,G,T,P} (2048-1  
04857599, default 104857599): +2G
```

```
Created a new partition 1 of type 'Linux' and of size  
2 GiB.
```

```
Command (m for help): w
```

```
The partition table has been altered.
```

```
Calling ioctl() to re-read partition table.
```

3. Did the kernel feel the changes? Display the content of /proc/partitions file? What did you notice? How to overcome that?

```
[root@server ~]# lsblk
NAME                                MAJ:MIN RM  SIZE RO TYPE MOUNTPOINTS
sr0                                11:0    1 10.3G  0  rom
nvme0n1                            259:0    0   50G  0  disk
├─nvme0n1p1                        259:1    0    1G  0  part /boot
├─nvme0n1p2                        259:2    0   49G  0  part
│   └─rhel_server-root
│       253:0    0   47G  0  lvm  /
│   └─rhel_server-swap
│       253:1    0    2G  0  lvm  [SWAP]
nvme0n2                            259:3    0   50G  0  disk
└─nvme0n2p1                        259:4    0    2G  0  part
[root@server ~]# cat /proc/partitions
major minor  #blocks  name

 259        0   52428800 nvme0n1
 259        1    1048576 nvme0n1p1
 259        2   51379200 nvme0n1p2
 259        3   52428800 nvme0n2
 259        4    2097152 nvme0n2p1
  11        0   10825920 sr0
 253        0   49262592 dm-0
 253        1    2113536 dm-1
[root@server ~]#
```

Yes feel the change and if else we can use partprobe

4. Make a new `ext4` file system on the new logical partition you just created.

```
[root@server ~]# mkfs.ext4 /dev/nvme0n2p1
mke2fs 1.46.5 (30-Dec-2021)
Creating filesystem with 524288 4k blocks and 131072 inodes
Filesystem UUID: f0f2d270-6d78-4611-b0a7-da45ab7442ee
Superblock backups stored on blocks:
    32768, 98304, 163840, 229376, 294912

Allocating group tables: done
Writing inode tables: done
Creating journal (16384 blocks): done
Writing superblocks and filesystem accounting information: done

[root@server ~]#
```

Bonus: Try creating the `ext4` filesystem with 2k blocks and one inode per every 4k(two blocks) of filesystem.

```
[root@server ~]# mkfs.ext4 -b 2048 -i 4096 /dev/nvme0n2p1
mke2fs 1.46.5 (30-Dec-2021)
/dev/nvme0n2p1 contains a ext4 file system
    created on Thu Nov 21 16:10:23 2024
Proceed anyway? (y,N) y
Creating filesystem with 1048576 2k blocks and 524288 inodes
Filesystem UUID: 22f5b5c1-25b8-493f-8645-ec5210a2df77
Superblock backups stored on blocks:
    16384, 49152, 81920, 114688, 147456, 409600, 442368, 802816

Allocating group tables: done
Writing inode tables: done
Creating journal (16384 blocks): done
Writing superblocks and filesystem accounting information: done

[root@server ~]#
```

5. Create a directory, name it /data.

```
[root@server ~]# mkdir /data
[root@server ~]# ls -ld/data
ls: invalid option -- '/'
Try 'ls --help' for more information.
[root@server ~]# ls -ld /data
drwxr-xr-x. 2 root root 6 Nov 21 16:15 /data
[root@server ~]#
```

6. Add a label to the new filesystem, name it data.

```
[root@server ~]# e2label /dev/nvme0n2p1 data
[root@server ~]#
```

7. Add a new entry to /etc/fstab for the new filesystem using the label you just create.

```
[root@server ~]# vi /etc/fstab
[root@server ~]#
LABEL=data /data ext4 defaults 0 2
```

8. Mount the new filesystem

```
[root@server ~]# mount -a
mount: (hint) your fstab has been modified, but systemd still uses
the old version; use 'systemctl daemon-reload' to reload.
[root@server ~]# mount /data
mount: /data: /dev/nvme0n2p1 already mounted on /data.
mount: (hint) your fstab has been modified, but systemd still uses
the old version; use 'systemctl daemon-reload' to reload.
[root@server ~]# lsblk
NAME                                MAJ:MIN RM  SIZE RO TYPE MOUNTPOINTS
sr0                                  11:0    1 10.3G  0 rom
nvme0n1                             259:0    0   50G  0 disk
├─nvme0n1p1                         259:1    0    1G  0 part /boot
└─nvme0n1p2                         259:2    0   49G  0 part
   ├─rhel_server-root              253:0    0   47G  0 lvm  /
   └─rhel_server-swap              253:1    0    2G  0 lvm  [SWAP]
nvme0n2                             259:3    0   50G  0 disk
└─nvme0n2p1                       259:4    0    2G  0 part /data
[root@server ~]#
```

9. Display your swap size.

10. Create a swap file of size 512MB.

11. Add the swap file to the virtual memory of the system.

12. Display the swap size.

13. Use the `fdisk` command to create 2 Linux LVM (0x8e) partitions using "unpartitioned" space on your hard disk. These partitions should all be the same size; to speed up the lab, do not make them larger than 300 MB each. Make sure to write the changes to disk by using the `w` command to exit the `fdisk` utility. Run the `partprobe` command after exiting the `fdisk` utility.

14. Initialize your Linux LVM partitions as physical volumes with the `pvcreate` command. You can use the `pvdisplay` command to verify that the partitions have been initialized as physical volumes.

15. Using only one of your physical volumes, create a volume group called `test0`. Use the `vgdisplay` command to verify that the volume group was created.

16. Create a small logical volume (LV) called `data` that uses about 30 percent of the available space of the `test0` volume group. Look for VG Size and Free PE/Size in the output of the `vgdisplay` command to assist you with this. Use the `lvdisplay` command to verify your work.

17. Create an `xfs` filesystem on your new LV.

18. Make a new directory called `/data` and then mount the new LV under the `/data` directory. Create a "large file" in this volume.

19. Enlarge the LV that you created in Sequence 1 (`/dev/test0/data`) by using approximately 25 percent of the remaining free space in the `test0` volume group. Then, enlarge the filesystem of the LV.

20. Verify that the file /data/bigfile still exists in the LV. Run the df command and check to verify that more free disk space is now available on the LV.
21. Use the remaining extents in the test0 volume group to create a second LV called docs.
22. Run the vgdisplay command to verify that there are no free extents left in the test0 volume group.
23. Create an xfs filesystem on the new LV, make a mount point called /docs and mount the docs LV using this mount point.
24. Add all of the remaining unused physical volumes that you created in Sequence 1 to the test0 volume group.
25. If you run vgdisplay again, there now should be free extents (provided by the new physical volumes) in the test0 volume group. Extend the docs LV and underlying filesystem to make use of all of the free extents of the test0 volume group. Verify your actions.